

FIRST 120 DAYS

Earn the roadmap before announcing it.

An entry plan for a Director, Applied ML/AI joining a broad global entrepreneur platform.

ENTRY PRINCIPLES

Start with customer movement

Connect every technical conversation to an entrepreneur decision, action, or outcome.

Read before prescribing

Respect existing teams, architecture, experiments, and constraints before introducing a new taxonomy.

Make authority explicit

Agent autonomy, customer control, Guide authority, and production ownership are design variables.

Use evidence to sequence

The three-year roadmap should reflect dependencies, reuse, economics, and proof - not executive enthusiasm alone.

DAYS 1-30 - READ THE RESOLUTION SYSTEM

| Workstream            | Actions                                                                                                               | Evidence produced                                             |
|-----------------------|-----------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------|
| Entrepreneur journeys | Map high-volume and high-friction moments from idea through launch, growth, commerce, service, and renewal.           | Journey/outcome map and candidate moments for personalization |
| AI/ML landscape       | Inventory models, agents, pipelines, data assets, experimentation tooling, evaluations, and production ownership.     | Capability and dependency map                                 |
| Authority and trust   | Review customer controls, Guide interventions, sensitive decisions, privacy/security boundaries, and regional policy. | Authority boundary register                                   |
| Operating reality     | Shadow engineering, product, analytics, Guides, security, privacy, finance, and incident/release routines.            | Constraint and decision-right map                             |

DAY-30 DECISIONS

- Which two entrepreneur moments deserve deeper discovery?
- What shared platform capabilities already exist and should be protected?
- Where is evidence weak, ownership ambiguous, or authority unclear?
- Which assumed constraints are real, and which are historical habits?

## DAYS 31-60 - DEFINE THE CONTRACTS

| Contract | Questions to settle                                                               | Output                             |
|----------|-----------------------------------------------------------------------------------|------------------------------------|
| Outcome  | What customer movement makes the work valuable?                                   | Outcome tree with measurable units |
| STATE    | Which context is required, permitted, fresh, and trustworthy?                     | Context/data contract              |
| NEXT     | What decision or action is the system optimizing?                                 | Next-best-move definition          |
| AUTH     | Who can draft, recommend, execute, approve, override, or recover?                 | Authority and escalation contract  |
| PROOF    | What evaluation, online evidence, trust signals, and economics earn continuation? | Evaluation and measurement plan    |
| TTL      | When must the state and decision be refreshed?                                    | Freshness and re-resolution policy |

## TECHNICAL CREDIBILITY REVIEW

### Architecture and pipeline

- Data ingestion and feature/context availability
- Training, evaluation, deployment, rollback, and monitoring
- Online inference, latency, cost, and failure modes
- Regional and product integration patterns

### Agentic behavior

- Tool access and permission boundaries
- Memory/context and state management
- Human review and Guide escalation
- Evaluation suites, traces, incident recovery, and learning

## DAYS 61-90 - RUN BOUNDED PROOFS

1. Select one lower-risk and one higher-leverage entrepreneur moment.
2. Establish current behavior and comparable baseline units.
3. Run controlled tests with explicit authority and rollback.
4. Instrument relevance, customer control, quality, economics, latency, operating load, and regional fit.
5. Review weekly with product, engineering, analytics, Guides, security/privacy, and finance.
6. Convert results into reusable components, documented limits, or a fast no.

DAYS 91-120 - COMMIT THE THREE-YEAR LOGIC

Roadmap architecture

- Customer moments sequenced by value, readiness, and dependency
- Shared platform capabilities and domain-owned experiences
- Evaluation, observability, governance, and regional foundations
- Explicit investment, hold, and stop decisions

Operating system

- Quarterly outcome and capability review
- Weekly experiment and production evidence review
- Incident, rollback, and learning capture
- Team, partner, hiring, and skill-development plan

ILLUSTRATIVE ROADMAP LANDMARKS

| Horizon     | Landmark                                                                                         | Evidence gate                                                      |
|-------------|--------------------------------------------------------------------------------------------------|--------------------------------------------------------------------|
| 0-6 months  | Two bounded customer-moment proofs; shared authority/evaluation contract                         | Observable momentum, customer control, reliable operations         |
| 6-12 months | Reusable context, experimentation, and observability components adopted across multiple journeys | Reuse, faster trustworthy experiments, acceptable cost/latency     |
| Year 2      | Cross-product next-best-move orchestration with regional controls and Guide integration          | Consistent standards with local fit; reduced friction and rework   |
| Year 3      | Adaptive entrepreneur platform learning across lifecycle moments                                 | Compounding customer and business outcomes without authority drift |

SUCCESS MEASURES FOR THE FIRST 120 DAYS

- A shared mandate and customer-outcome vocabulary is used by technical and business leaders.
- Two entrepreneur moments have explicit state, action, authority, proof, and freshness contracts.
- Current architecture and operating constraints are documented without blame.
- Technical and nontechnical stakeholders agree on scale/hold/stop evidence.
- The roadmap distinguishes shared platform work, domain ownership, and local adaptation.
- Technical credibility is earned through architecture review, direct tooling fluency, specialist trust, deliberate hiring and partnering, and visible learning.

WHAT RUSSELL WOULD NOT DO

Announce a universal agent architecture before reading the system; treat engagement as proof of customer value; centralize every AI decision; automate sensitive exceptions without authority design; or present a three-year roadmap whose dependencies and stop criteria are invisible.