

FIRST 120 DAYS

Earn the roadmap before announcing it.

An entry plan for a Director, Applied ML/AI joining a broad global entrepreneur platform.

ENTRY PRINCIPLES

Start with customer movement

Connect every technical conversation to an entrepreneur decision, action, or outcome.

Read before prescribing

Respect existing teams, architecture, experiments, and constraints before introducing a new taxonomy.

Make authority explicit

Agent autonomy, customer control, Guide authority, and production ownership are design variables.

Use evidence to sequence

The three-year roadmap should reflect dependencies, reuse, economics, and proof - not executive enthusiasm alone.

DAYS 1-30 - READ THE RESOLUTION SYSTEM

Workstream	Actions	Evidence produced
Entrepreneur journeys	Map high-volume and high-friction moments from idea through launch, growth, commerce, service, and renewal.	Journey/outcome map and candidate moments for personalization
AI/ML landscape	Inventory models, agents, pipelines, data assets, experimentation tooling, evaluations, and production ownership.	Capability and dependency map
Authority and trust	Review customer controls, Guide interventions, sensitive decisions, privacy/security boundaries, and regional policy.	Authority boundary register
Operating reality	Shadow engineering, product, analytics, Guides, security, privacy, finance, and incident/release routines.	Constraint and decision-right map

DAY-30 DECISIONS

- Which two entrepreneur moments deserve deeper discovery?
- What shared platform capabilities already exist and should be protected?
- Where is evidence weak, ownership ambiguous, or authority unclear?
- Which assumed constraints are real, and which are historical habits?

DAYS 31-60 - DEFINE THE CONTRACTS

Contract	Questions to settle	Output
Outcome	What customer movement makes the work valuable?	Outcome tree with measurable units
STATE	Which context is required, permitted, fresh, and trustworthy?	Context/data contract
NEXT	What decision or action is the system optimizing?	Next-best-move definition
AUTH	Who can draft, recommend, execute, approve, override, or recover?	Authority and escalation contract
PROOF	What evaluation, online evidence, trust signals, and economics earn continuation?	Evaluation and measurement plan
TTL	When must the state and decision be refreshed?	Freshness and re-resolution policy

TECHNICAL CREDIBILITY REVIEW

Architecture and pipeline

- Data ingestion and feature/context availability
- Training, evaluation, deployment, rollback, and monitoring
- Online inference, latency, cost, and failure modes
- Regional and product integration patterns

Agentic behavior

- Tool access and permission boundaries
- Memory/context and state management
- Human review and Guide escalation
- Evaluation suites, traces, incident recovery, and learning

DAYS 61-90 - RUN BOUNDED PROOFS

1. Select one lower-risk and one higher-leverage entrepreneur moment.
2. Establish current behavior and comparable baseline units.
3. Run controlled tests with explicit authority and rollback.
4. Instrument relevance, customer control, quality, economics, latency, operating load, and regional fit.
5. Review weekly with product, engineering, analytics, Guides, security/privacy, and finance.
6. Convert results into reusable components, documented limits, or a fast no.

DAYS 91-120 - COMMIT THE THREE-YEAR LOGIC

Roadmap architecture

- Customer moments sequenced by value, readiness, and dependency
- Shared platform capabilities and domain-owned experiences
- Evaluation, observability, governance, and regional foundations
- Explicit investment, hold, and stop decisions

Operating system

- Quarterly outcome and capability review
- Weekly experiment and production evidence review
- Incident, rollback, and learning capture
- Team, partner, hiring, and skill-development plan

ILLUSTRATIVE ROADMAP LANDMARKS

Horizon	Landmark	Evidence gate
0-6 months	Two bounded customer-moment proofs; shared authority/evaluation contract	Observable momentum, customer control, reliable operations
6-12 months	Reusable context, experimentation, and observability components adopted across multiple journeys	Reuse, faster trustworthy experiments, acceptable cost/latency
Year 2	Cross-product next-best-move orchestration with regional controls and Guide integration	Consistent standards with local fit; reduced friction and rework
Year 3	Adaptive entrepreneur platform learning across lifecycle moments	Compounding customer and business outcomes without authority drift

SUCCESS MEASURES FOR THE FIRST 120 DAYS

- A shared mandate and customer-outcome vocabulary is used by technical and business leaders.
- Two entrepreneur moments have explicit state, action, authority, proof, and freshness contracts.
- Current architecture and operating constraints are documented without blame.
- Technical and nontechnical stakeholders agree on scale/hold/stop evidence.
- The roadmap distinguishes shared platform work, domain ownership, and local adaptation.
- Known candidate capability gaps are addressed through technical review, hiring, partnering, and deliberate learning.

WHAT RUSSELL WOULD NOT DO

Announce a universal agent architecture before reading the system; treat engagement as proof of customer value; centralize every AI decision; automate sensitive exceptions without authority design; or present a three-year roadmap whose dependencies and stop criteria are invisible.